

hw3

Jiahao Tian

2023-02-27

Question 1

a) Reversible Jumps MCMC algorithm

$$\begin{aligned} p(y|\beta, \sigma^2, \gamma)p(\beta|\sigma^2, \gamma) &= N(y; X_\gamma\beta, \sigma^2 I)N(\beta; \frac{\sigma^2}{g}(X_\gamma^\top X_\gamma)^{-1}) \\ &= (2\pi\sigma^2)^{-n/2} \exp\left(-\frac{1}{2\sigma^2}(y - X_\gamma\beta)^\top (y - X_\gamma\beta)\right) \exp\left(-\frac{g}{2\sigma^2}\beta^\top X_\gamma^\top X_\gamma\beta\right) \\ &\quad - \frac{1}{2\sigma^2}(y - X_\gamma\beta)^\top (y - X_\gamma\beta) - \frac{g}{2\sigma^2}\beta^\top X_\gamma^\top X_\gamma\beta \\ &= -\frac{1}{2}\left[y^\top y/\sigma^2 - 2\beta^\top X_\gamma y/\sigma^2 + \beta^\top X_\gamma^\top X_\gamma\beta/\sigma^2 + g\beta^\top X_\gamma^\top X_\gamma\beta/\sigma^2\right] \\ &= -\frac{1}{2}\left[y^\top y/\sigma^2 + (\beta - m)M^{-1}(\beta - m) + m^\top M^{-1}m\right] \end{aligned}$$

where $M = \sigma^2(X_\gamma^\top X_\gamma)^{-1}/(1 + g)$ and $m = MX^\top y/\sigma^2$

$$\text{So } p(y|\beta, \sigma^2, \gamma)p(\beta|\sigma^2, \gamma) = (2\pi\sigma^2)^{-n/2} \exp\left(-\frac{1}{2}\left[y^\top y/\sigma^2 + (\beta - m)M^{-1}(\beta - m) + m^\top M^{-1}m\right]\right)$$

- Now need to integrate it out.

$$\begin{aligned}\int p(y|\beta, \sigma^2, \gamma) p(\beta|\sigma^2, \gamma) d\beta &= (2\pi\sigma^2)^{-n/2} \exp(\frac{1}{2\sigma^2} y^\top y) \exp(\frac{1}{2} m^\top M^{-1} m) |2\pi M|^{1/2} \\ &= (2\pi)^{(-n+q_\gamma)/2} (\sigma^2)^{(-n+q_\gamma)/2} |(X_\gamma^\top X_\gamma)|^{-1/2} (\frac{1}{g+1})^{q_\gamma/2} \exp(-\frac{1}{2\sigma^2} y^\top y) \exp(\frac{1}{2} m^\top M^{-1} m)\end{aligned}$$

$$\begin{aligned}\exp(-\frac{1}{2\sigma^2} y^\top y) \exp(\frac{1}{2} m^\top M^{-1} m) &= \exp\left(-\frac{1}{2\sigma^2} (y^\top y - y^\top X_\gamma (X_\gamma^\top X_\gamma)^{-1} X_\gamma^\top y / (1+g))\right) \\ &= \exp\left(-\frac{1}{2\sigma^2} y^\top \left(I - \frac{1}{1+g} X_\gamma (X_\gamma^\top X_\gamma)^{-1} X_\gamma^\top\right) y\right) \\ &= \exp\left(-\frac{1}{2\sigma^2} RSS\right)\end{aligned}$$

$$\begin{aligned}&= (2\pi)^{-(n+q_\gamma)/2} (\frac{1}{g+1})^{q_\gamma/2} |(X_\gamma^\top X_\gamma)|^{-1/2} \int (\sigma^2)^{(-n+q_\gamma)/2} \exp(-\frac{1}{2\sigma^2} RSS) p(\sigma^2) d\sigma^2 \\ &= (2\pi)^{(-n+q_\gamma)/2} (\frac{1}{g+1})^{q_\gamma/2} |(X_\gamma^\top X_\gamma)|^{-1/2} \frac{\Gamma(\frac{n}{2} - 1)}{(RSS/2)^{n/2}}\end{aligned}$$

which follows $\sim \Gamma^{-1}(n/2 - 1, RSS/2)$

- sample from $p(\gamma|Y)$.

$$\begin{aligned}p(\gamma|y) &\propto p(y|\gamma) p(\gamma) \\ &= (2\pi)^{-(n+q_\gamma)/2} (\frac{1}{g+1})^{q_\gamma/2} |(X_\gamma^\top X_\gamma)|^{-1/2} \frac{\Gamma(\frac{n}{2} - 1)}{(RSS/2)^{n/2}} p(\gamma)\end{aligned}$$

$$\begin{aligned}p(\sigma^2, \beta|\gamma, \sigma^2) &\propto p(y|\gamma, \beta_\gamma, \sigma^2) p(\beta_\gamma|\sigma^2, \gamma) p(\sigma^2|\gamma) \\ &= N(m, M) \times IG(n/2 - 1, RSS/2)\end{aligned}$$

$$\begin{aligned}p(\beta_\gamma, \gamma, \sigma^2|y) &\propto p(y|\beta_\gamma, \sigma^2, \gamma) p(\beta_\gamma, \sigma^2, \gamma) \\ &= p(y|\beta_\gamma, \sigma^2, \gamma) p(\beta_\gamma, \sigma^2|\gamma) p(\gamma) \\ &= N(m, M) \times IG(n/2 - 1, RSS/2) \times 2^{-q_\gamma}\end{aligned}$$

```
# Define the function to calculate the posterior probability
posterior_prob = function(x, y, gamma, alpha, beta, sigma2_0, tau2_0) {
  n = length(y)
  p_gamma = length(gamma)
  if (p_gamma == 0) {
    return(0)
  }
  else {
    x_gamma = x[, gamma, drop = FALSE]
    tau2 = tau2_0[gamma]
    sigma2 = sigma2_0
    beta_gamma = MASS::ginv(x_gamma) %*% y
    log_lik = -0.5 * (n * log(sigma2)
      + sum((y - x_gamma %*% beta_gamma)^2) / sigma2)
    log_prior = sum(dnorm(beta_gamma, mean = 0, sd = sqrt(tau2), log = TRUE)) +
      sum(dgamma(tau2, shape = alpha, rate = beta, log = TRUE))
  }
}
```

```

    return(log_lik + log_prior)
  }
}

```

```

inclusion_prob = function(x, y, n_iter, burn, alpha, beta, sigma2_0, tau2_0) {

  # Initialize variables
  n = ncol(x)
  gamma = numeric(n)
  gamma_old = gamma
  posterior_prob_old = -Inf
  posterior_prob_vec = numeric(n_iter - burn)
  gamma_mat = matrix(0, ncol = n, nrow = n_iter - burn)

  # Run the MCMC algorithm
  for (i in 1:n_iter) {
    # Propose a move
    if (runif(1) < 0.5) {
      # Add a variable
      gamma_prop = sample(n, sum(gamma == 0), replace = FALSE)
      gamma_prop = sort(c(gamma[gamma != 0], gamma_prop))
    }
    else {
      # Delete a variable (only if gamma is not empty)
      if (sum(gamma != 0) > 0) {
        gamma_prop = gamma[-sample(which(gamma != 0), 1)]
      }
      else {
        gamma_prop = gamma
      }
    }

    # Calculate acceptance probability
    posterior_prob_prop = posterior_prob(x, y, gamma_prop, alpha, beta, sigma2_0, tau2_0)
    alpha_prop = min(1, exp(posterior_prob_prop - posterior_prob_old))
    # Accept or reject move
    if (runif(1) < alpha_prop) {
      gamma = gamma_prop
      posterior_prob_old = posterior_prob_prop
    }

    # Save current state if past burn-in period
    if (i > burn) {
      gamma_mat[i - burn, gamma] = 1
      posterior_prob_vec[i - burn] = posterior_prob_old
    }
  }

  # Calculate inclusion probabilities
  inclusion_probs = colMeans(gamma_mat)
}

```

```

set.seed(888)
data(diabetes)

```

```

x = cbind(1, diabetes$x)
x = as.matrix(x)
y = diabetes$y

# Define prior parameters
alpha = 0.5
beta = 0.5
sigma2_0 = 100
tau2_0 = rep(10, ncol(x))

# Set number of iterations and burn
n_iter = 10000
burn = 5000

ip = inclusion_prob(x, y, n_iter, burn, alpha, beta, sigma2_0, tau2_0)
print(ip)

## [1] 0.8236 0.0000 0.0000 0.8166 0.8224 0.0000 0.8192 0.8220 0.8188 0.8218
## [11] 0.8232

```

- I tried my best...

b)

- Compare results in (a) with results obtained sampling directly from $p(\gamma|Y)$. Implement this without the need for RJ Monte Carlo methods.

```

# Define the function
p_x = function(y, x, g = 0.001) {

  # Initialize variables
  n = nrow(x)
  p = ncol(x)
  if(is.null(p)) {
    X = as.matrix(x)
    n = dim(x)[1]
    p = dim(x)[2]
  }
  H = 0
  if(p > 1) {
    H = 1 / (g + 1) * x %*% solve(t(x) %*% x) %*% t(x)
  }
  RSS = t(y) %*% (diag(n) - H) %*% y
  return((-n + p) / 2 * log(2 * pi) -
         p / 2 * log(g + 1) - 1 / 2 * log(det(t(x) %*% x)) -
         lgamma(n / 2 - 1) - n / 2 * log(RSS / 2))
}

```

```

inclusion_prob2 = function(x, y, n_iter) {

  # Initialize variables

```

```

p = ncol(x)
z = rep(1, p)
p_proposed = p_x(y, x[, z == 1])
p_current = p_proposed
Z = matrix(0, n_iter, ncol(x))

for(i in 1:n_iter) {
  for(j in sample(2:p)) {
    z.propose = z
    z.propose[j] = 1 - z.propose[j]

    # log-likelihood of proposed
    p_proposed = p_x(y, x[, z.propose == 1])

    r = (p_proposed - p_current) * (-1)^(z.propose[j] == 0)

    # sample from bernoulli
    z[j] = rbernoulli(1, 1 / (1 + exp(-r)))

    if(z[j] == z.propose[j]) {
      p_current = p_proposed
    }
  }
  Z[i, ] = z
}
inclusion_probs2 = colMeans(Z[1:(1 * n_iter), ])
}

```

```

set.seed(999)
y = diabetes$y
x = cbind(1, diabetes$x)
x = as.matrix(x)

ip2 = inclusion_prob2(x = x, y = y, n_iter = 5000)

```

```
## Warning: 'rbernoulli()' was deprecated in purrr 1.0.0.
```

```
print(ip2)
```

```
## [1] 1.0000 0.7324 0.9998 1.0000 1.0000 0.9884 0.9768 0.9104 0.9248 1.0000
## [11] 0.8374
```