

hw4

Jiahao Tian

2023-03-13

Problem: Consider Bayesian inference for data arising from counting processes. Specifically, given a set of covariates $X_i \in R^p$, and an associated set of regression coefficients $\beta \in R^p$, assume:

$$\begin{aligned} Y_i | \lambda_i &\sim \text{Poi}(e^{\lambda_i}) \\ \lambda_i | \beta, \tau &\sim N(X_i' \beta, \tau) \\ \beta &\sim N(0, (X'X)^{-1}g) \\ \tau &\sim IG(1, 1) \end{aligned}$$

a) Describe an MCMC strategy aimed at obtaining samples from the posterior distribution: $p(\beta, \tau | Y)$.

- By Gibbs sampling, we first need to find out the full conditional posterior distribution for λ, β, τ .
- Full conditional distributions have the same starting point: the full joint posterior distribution. Thus, the process of finding full conditional distributions is the same as finding the posterior distribution of each parameter.
- From the joint posterior distribution $p(\beta, \tau | Y)$, we can get:

a. Full conditional posterior distribution for β

$$\begin{aligned} p(\beta | \lambda_i, \tau) &= p(\lambda_i | \beta) P(\beta) \\ &\propto \prod \exp\left(-\frac{(\lambda_i - X_i \beta)^2}{2\tau}\right) \exp\left(-\frac{1}{2g} \beta' (X'X) \beta\right) \\ &= \exp\left(-\frac{1}{2\tau} (\beta' X' X \beta - 2\beta X' \lambda_i) - \frac{1}{2g} \beta' (X'X) \beta\right) \\ &= \exp\left(-\frac{1}{2} \left(\frac{\beta' X' X \beta - 2\beta X' \lambda_i}{\tau} - \beta' (X'X) g^{-1} \beta \right)\right) \\ &= \exp\left(-\frac{1}{2} \left(\frac{\beta' (X'X + (X'X)g^{-1}) \beta}{\tau} - 2\beta X' \lambda_i \right)\right) \\ &= \exp\left(\frac{1}{2} \left[\beta' X' X \left(\frac{1}{g} + \frac{1}{\tau} \right) \beta - 2 \frac{1}{g} \beta' X' \lambda_i \right]\right) \\ \text{which follows } &\sim N\left(\frac{X' \lambda_i \tau}{(X'X)\tau + (X'X)g}, \frac{g\tau}{(X'X)\tau + (X'X)g}\right) \end{aligned}$$

b. Full conditional posterior distribution for τ

$$\begin{aligned}
p(\tau|\lambda_i, \beta) &= p(\lambda_i|\tau)P(\tau) \\
&\propto \prod \exp\left(-\frac{(\lambda_i - X_i\beta)^2}{2\tau}\right) \tau^{-\frac{n}{2}} \tau^{-1-1} \exp\left(-\frac{1}{\tau}\right) \\
&= \tau^{-1-1} \tau^{-\frac{n}{2}} \exp\left(-\frac{1}{2\tau}(\lambda_i - X_i\beta)'(\lambda_i - X_i\beta) - \frac{1}{\tau}\right) \\
&= \tau^{-1-1} \tau^{-\frac{n}{2}} \exp\left(-\frac{1}{\tau}\left(\frac{(\lambda_i - X_i\beta)'(\lambda_i - X_i\beta)}{2} + 1\right)\right) \\
&\text{which follows } \sim IG\left(1 + \frac{n}{2}, 1 + \frac{1}{2}((\lambda_i - X_i\beta)'(\lambda_i - X_i\beta))\right)
\end{aligned}$$

c. Full conditional posterior distribution for λ

$$\begin{aligned}
p(\lambda_i | \beta, \tau, Y_i) &= p(Y_i | \lambda_i)p(\lambda_i | \beta, \tau) \\
&= \prod_{i=1}^n \frac{\exp(e^{\beta})e^{\beta Y_i}}{Y_i!} \cdot \frac{1}{\sqrt{2\pi\tau}} \exp\left(-\frac{(\lambda_i - X_i\beta)^2}{2\tau}\right) \\
&\text{because it is poisson distribution } e^{\lambda_i} \text{ needs to be greater than 0} \\
&\propto I_{e^{\lambda_i} \geq 0} \cdot \exp\left(-\frac{(\lambda_i - X_i\beta)^2}{2\tau}\right) \exp(\lambda_i Y_i) \exp(-e^{\lambda_i})
\end{aligned}$$

- Then we can take turns sampling these distributions like so:

- Using $\lambda_{i-1}, \beta_{i-1}$, draw τ_i from $p(\tau|\beta = \beta_{i-1}, \lambda = \lambda_{i-1})$.
- Using $\lambda_{i-1}, \tau_{i-1}$, draw β_i from $p(\beta|\tau = \tau_{i-1}, \lambda = \lambda_{i-1})$.
- Using β_{i-1}, τ_{i-1} , draw λ_i from $p(\lambda_i|\tau = \tau_{i-1}, \beta = \beta_{i-1})$.
- Together, steps 1, 2 and 3 complete one cycle of the Gibbs sampler and produce the draw for $(\beta_i, \tau_i, \lambda_i)$ in one iteration of a MCMC sampler.

b) Describe an HMC strategy aimed at obtaining samples from the posterior distribution: $p(\beta, \tau|Y)$.

- We first need to find out $\nabla U(x)$, $U(x) = -\log\pi(x)$.
- a. for β .

$$\begin{aligned}
p(\beta|\lambda_i, \tau) &= p(\lambda_i|\beta)P(\beta) \\
&\propto \exp\left(-\frac{1}{2}\left(\frac{\beta'(X'X + (X'X)g^{-1})\beta}{\tau} - 2\beta X' \lambda_i\right)\right) \\
&\text{which follows } \sim N\left(\frac{X' \lambda_i \tau}{(X'X)\tau + (X'X)g}, \frac{g\tau}{(X'X)\tau + (X'X)g}\right) \\
\nabla(-\log p(\beta|\lambda_i, \tau)) &\propto \beta\left(\frac{(X'X + (X'X)g^{-1})}{\tau} - X' \lambda_i\right)
\end{aligned}$$

- b. for τ .

$$\begin{aligned}
p(\tau|\lambda_i, \beta) &= p(\lambda_i|\tau)P(\tau) \\
&\propto \tau^{-1-1}\tau^{-\frac{n}{2}} \exp\left(-\frac{1}{\tau}\left(\frac{(\lambda_i - X_i\beta)'(\lambda_i - X_i\beta)}{2} + 1\right)\right) \\
&\text{which follows } \sim IG\left(1 + \frac{n}{2}, 1 + \frac{1}{2}((\lambda_i - X_i\beta)'(\lambda_i - X_i\beta))\right) \\
\nabla(-\log p(\tau|\lambda_i, \beta)) &\propto \frac{\frac{n}{2} + 2}{\tau} - \frac{(\lambda_i - X_i\beta)'(\lambda_i - X_i\beta)}{2\tau^2} - \frac{1}{\tau^2}
\end{aligned}$$

- c. for λ .

$$\begin{aligned}
p(\lambda_i | \beta, \tau, Y_i) &= p(Y_i | \lambda_i)p(\lambda_i | \beta, \tau) \\
&\propto \exp\left(e^{\lambda_i} + \lambda_i Y_i - \frac{(\lambda_i - X_i\beta)^2}{2\tau}\right) \\
\nabla(-\log p(\lambda_i|\tau, \beta, Y_i)) &\propto \frac{2(\lambda_i - X_i\beta)}{2\tau} - e^{\lambda_i} + Y_i
\end{aligned}$$

- After finds out $\nabla U(x)$ for each parameters, we can use leapfrog then implement HMC sampler in MCMC.

c) Test algorithms in (a) and (b) on the following data:

```

y = c(18,17,15,20,10,20,25,13,12)
x1 = gl(3,1,9)
x2 = gl(3,3)
dat = data.frame(y, x1, x2)

y = dat$y
X = cbind(rep(1, length(y)), dat$x1, dat$x2)

```

From part a):

```

update_beta = function (g, X, lamb, tau) {
  xtx = solve(t(X) %*% X)
  bmean = g / (g + tau) * xtx %*% t(X) %*% lamb
  bvar = tau * g / (g + tau) * xtx
  rmvnorm(1, mean = bmean, sigma = bvar)
}

update_tau = function (n, X, beta, lamb) {
  xb = X %*% t(beta)
  shape_1 = n / 2 + 1
  rate_1 = 1 + 1 / 2 * t(lamb - xb) %*% (lamb - xb)
  1 / rgamma(1, shape = shape_1, rate = rate_1 / 2)
}

```

```
## log lambda
update_lg_lamb = function(lambi, Xi, yi, beta, tau) {
  return(-1 / (2 * tau) * (lambi - Xi %*% t(beta))^2 + lambi * yi - exp(lambi))
}
```

```
gibs = function (n_iter, burn, X, y, cand_sd) {

  n = length(y)
  p = ncol(X)
  g = n

  ## initialize
  lamb_now = rep(1, n)
  beta_now = rep(1, p)
  tau_now = 1

  ## save chain
  beta_out = matrix(NA, nrow = n_iter, ncol = p)
  lamb_out = matrix(NA, nrow = n_iter, ncol = n)
  tau_out = rep(NA, n_iter)

  ## Gibbs sampler
  for(i in 1:n_iter) {

    # beta
    beta_now = update_beta(g = g, X = X, lamb = lamb_now, tau = tau_now)

    # tau
    tau_now = update_tau(n = n, X = X, beta = beta_now, lamb = lamb_now)

    # Random-Walk Metropolis-Hastings for lambda
    for(idx in 1:n) {

      ## step 1, initialize
      lambi = lamb_now[idx]
      Xi = X[idx, ]
      yi = y[idx]

      ## step 2, iterate
      lamb_cand = rnorm(1, lambi, cand_sd) # draw a candidate
      lamb_0 = update_lg_lamb(lambi = lambi,
                             Xi = Xi,
                             yi = yi,
                             beta = beta_now,
                             tau = tau_now)

      # evaluate with the candidate
      lamb_1 = update_lg_lamb(lambi = lamb_cand,
                             Xi = Xi,
                             yi = yi,
                             beta = beta_now,
                             tau = tau_now)

      ratio = lamb_1 - lamb_0 # log of acceptance ratio
    }
  }
}
```

```

    if (log(runif(1)) < ratio) { # accept the candidate
      lamb_now[idx] = lamb_cand
    }
  }
  ## collect results
  beta_out[i,] = beta_now
  lamb_out[i,] = lamb_now
  tau_out[i] = tau_now
}

cbind(beta = beta_out, tau = tau_out, lamb = lamb_out)
}

```

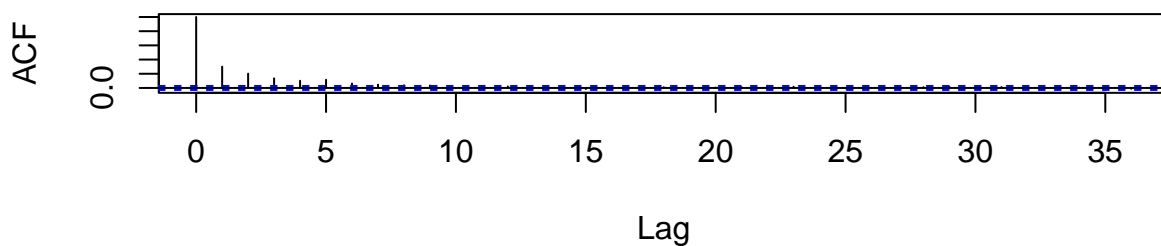
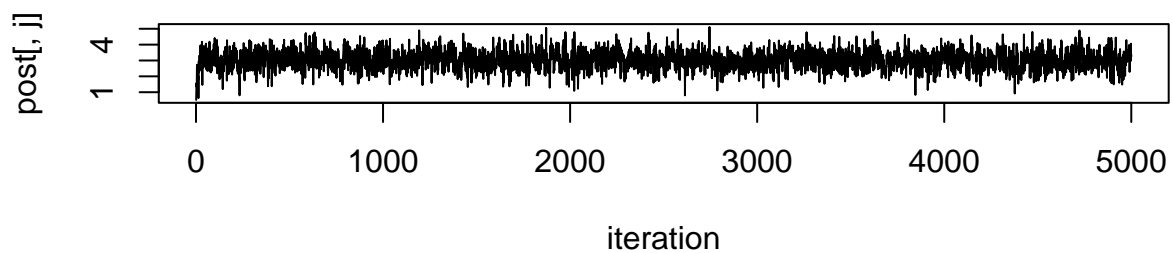
```
set.seed(11111)
```

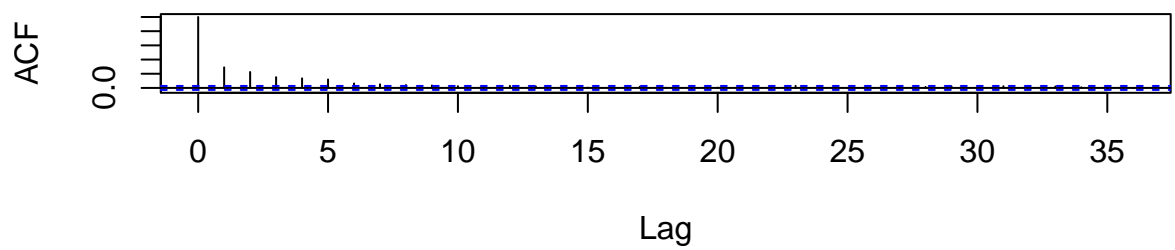
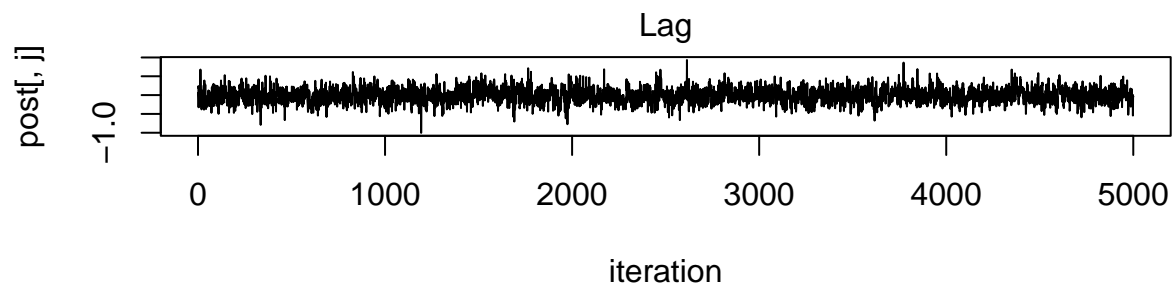
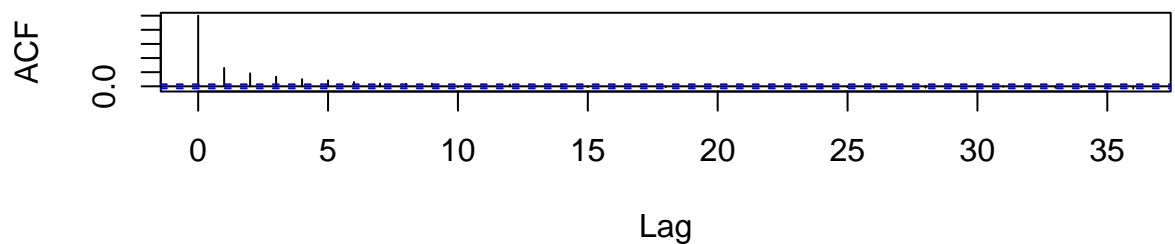
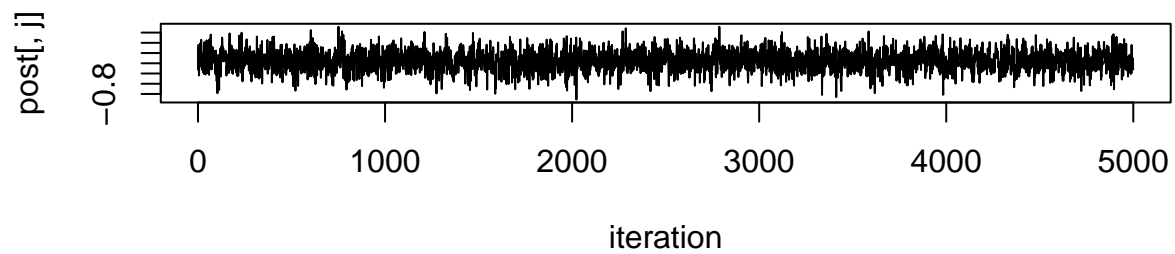
```
post = gibs(n_iter = 5000, burn = 1000, X = X, y = y, cand_sd = 0.5)
```

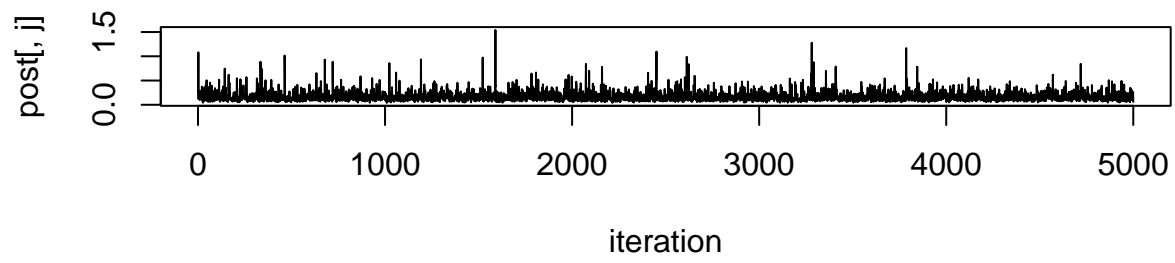
```

par(mfrow = c(2,1))
for(j in 1:13){
  matplot(c(1:5000), post[,j], type = "l", xlab = "iteration")
  acf(post[, j], main = colnames(post)[j]) }

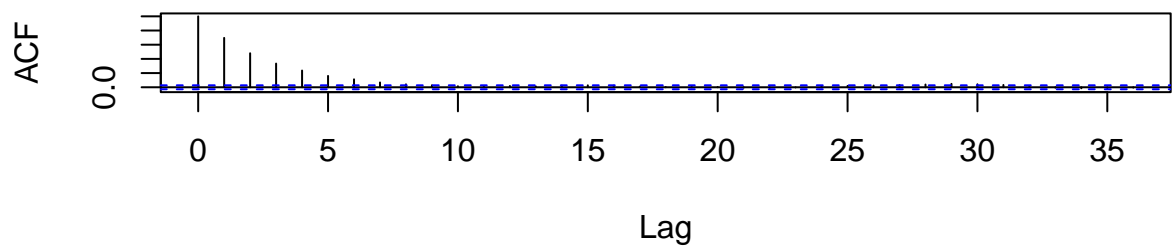
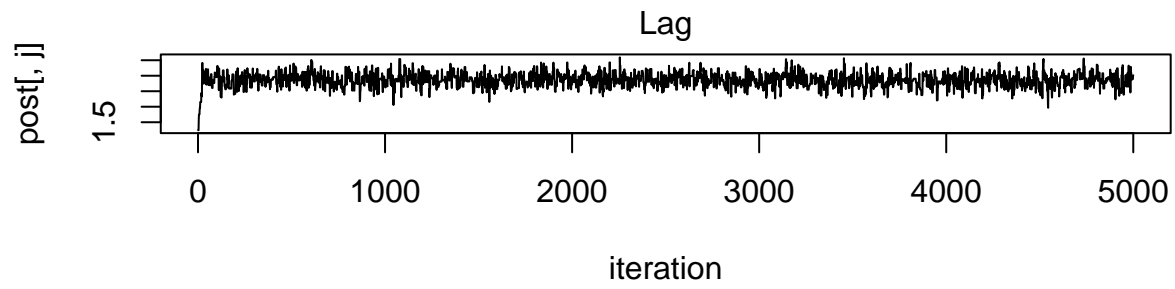
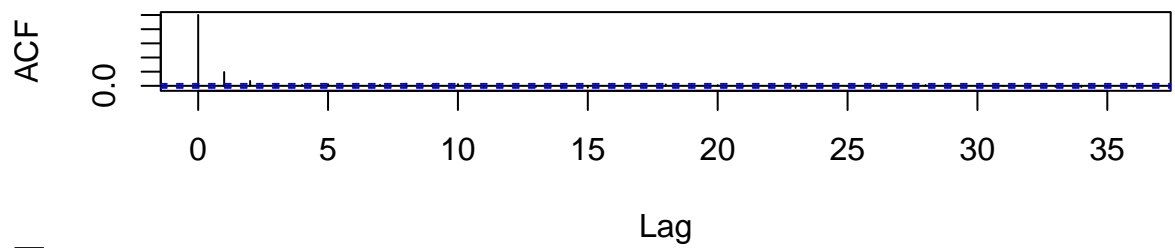
```

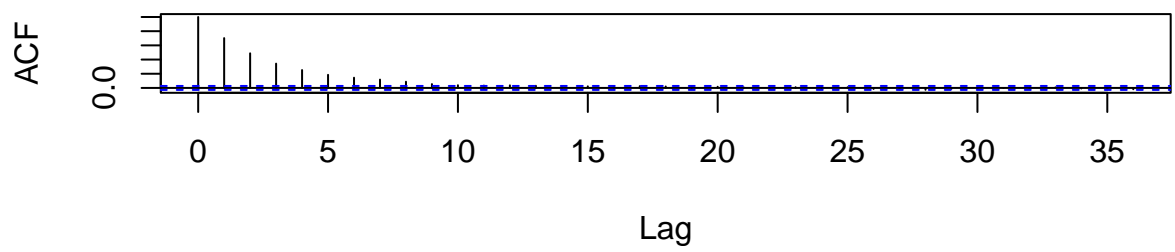
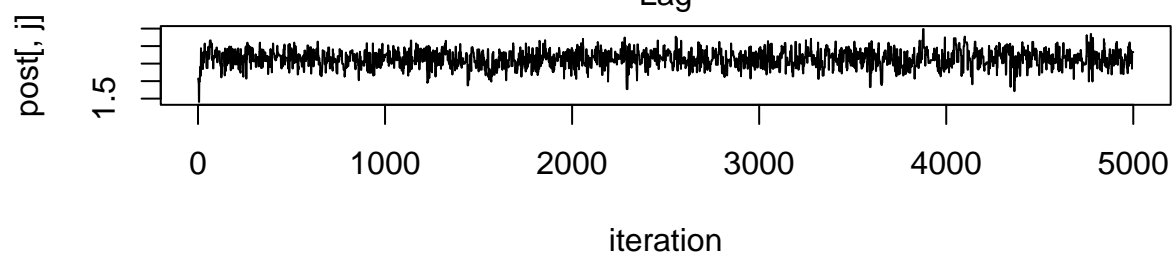
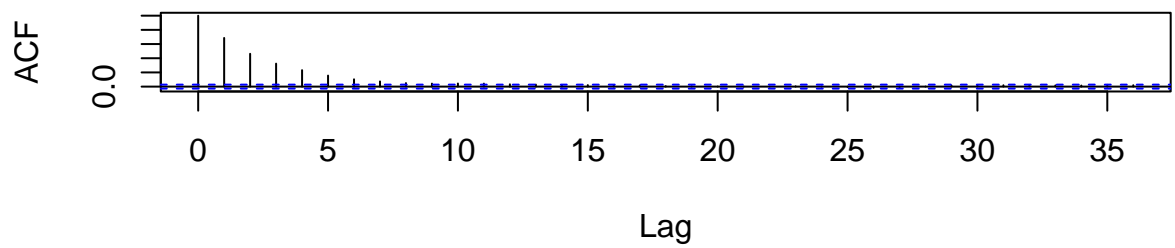
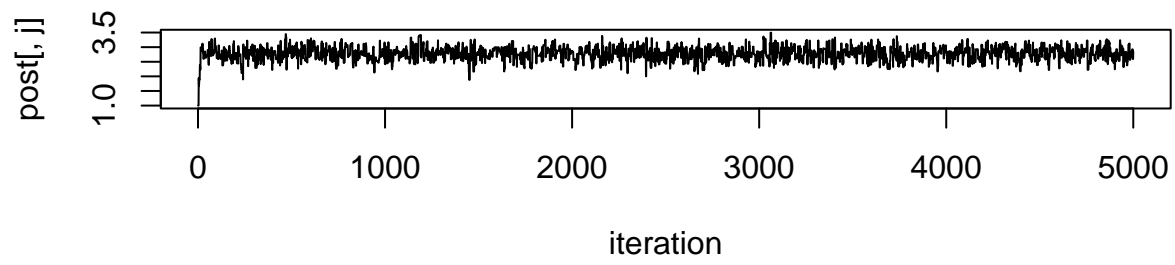


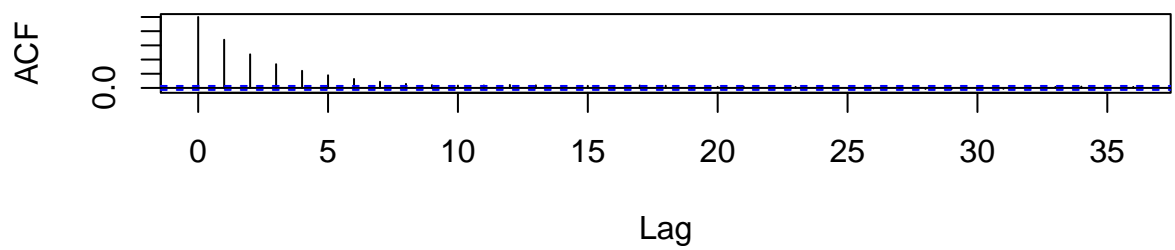
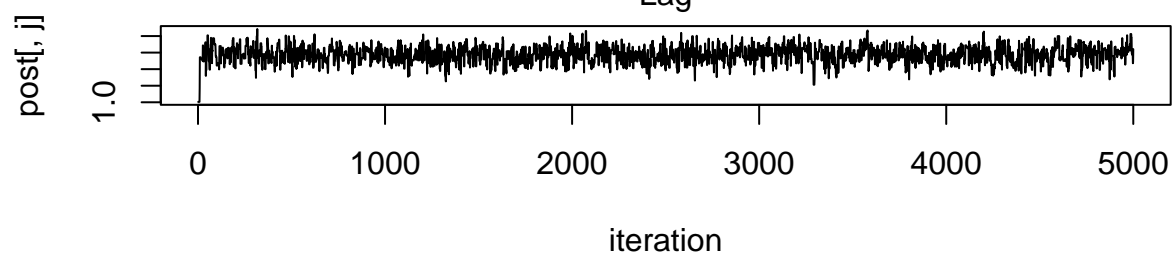
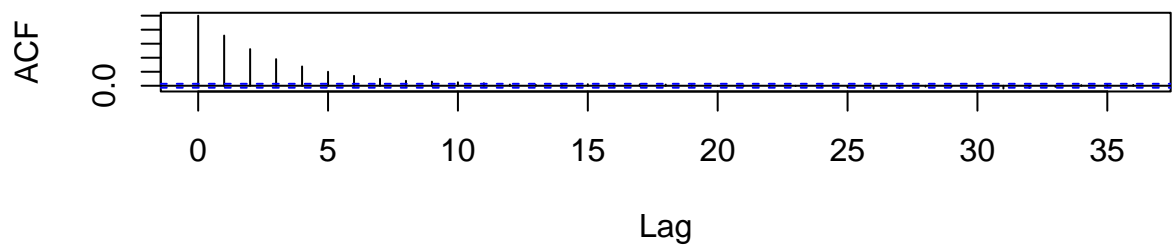
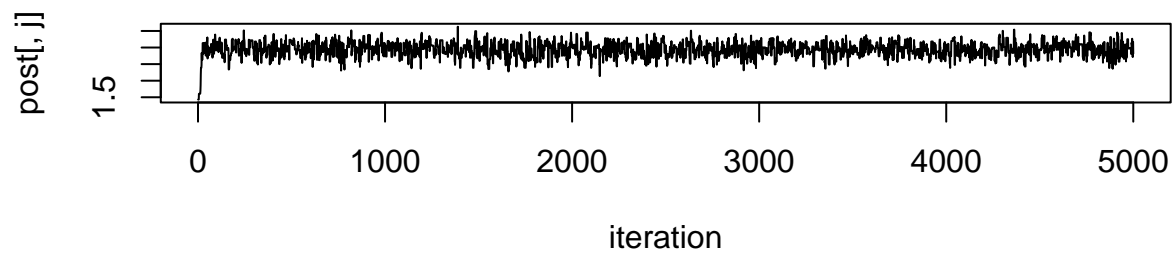


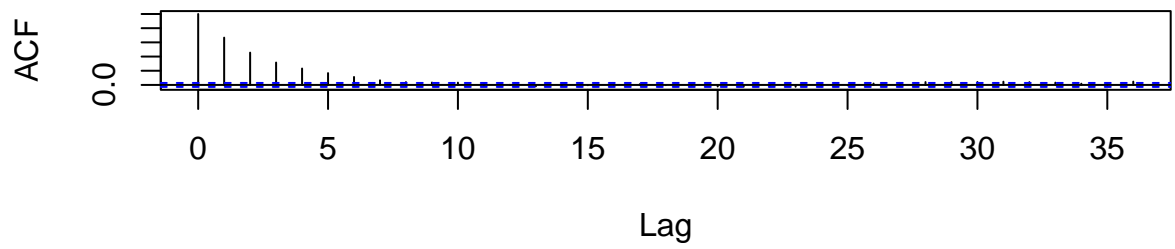
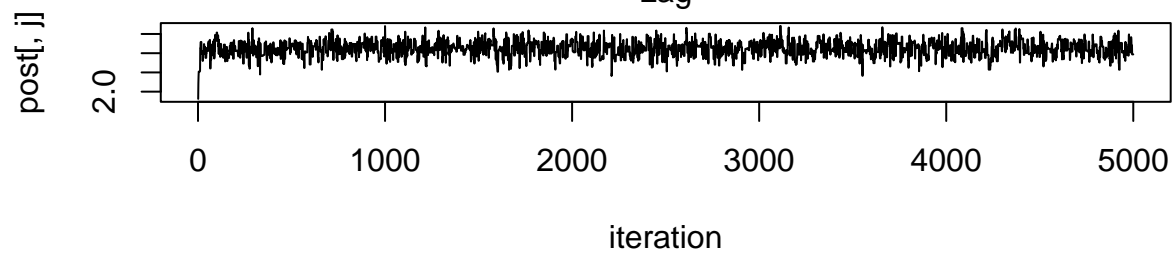
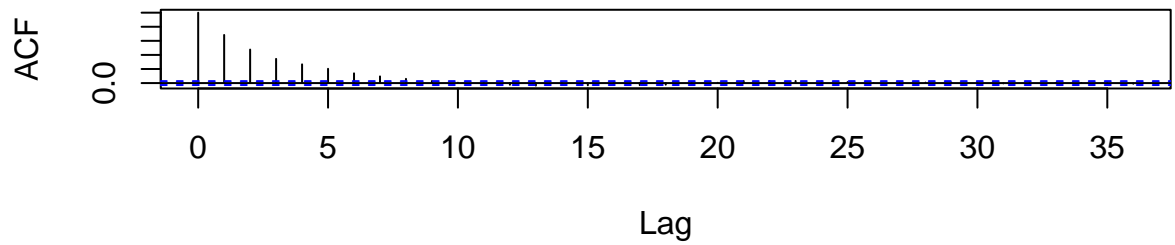
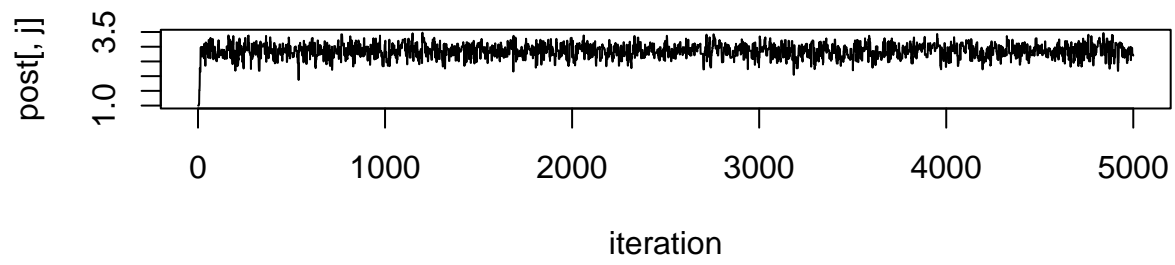


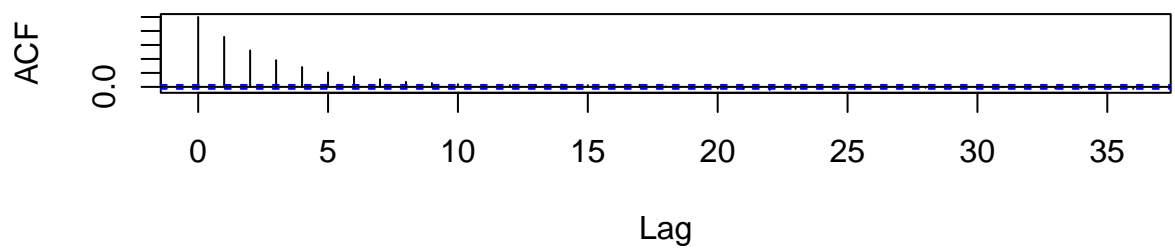
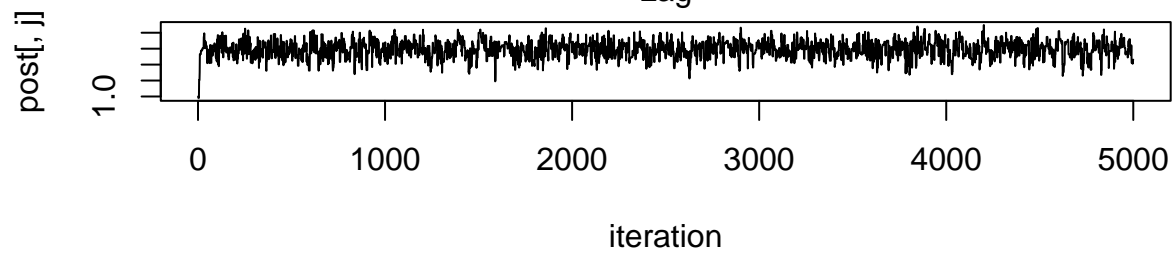
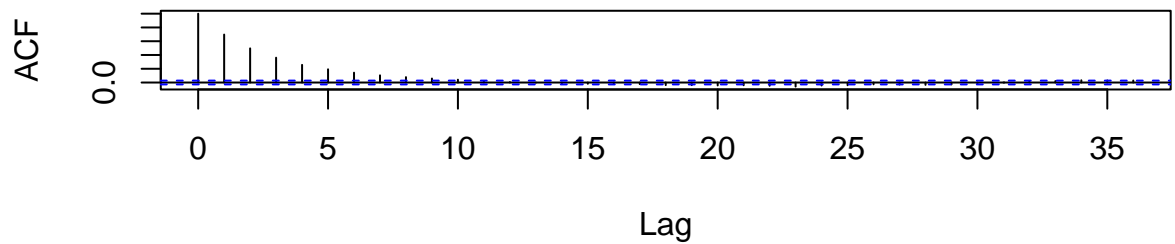
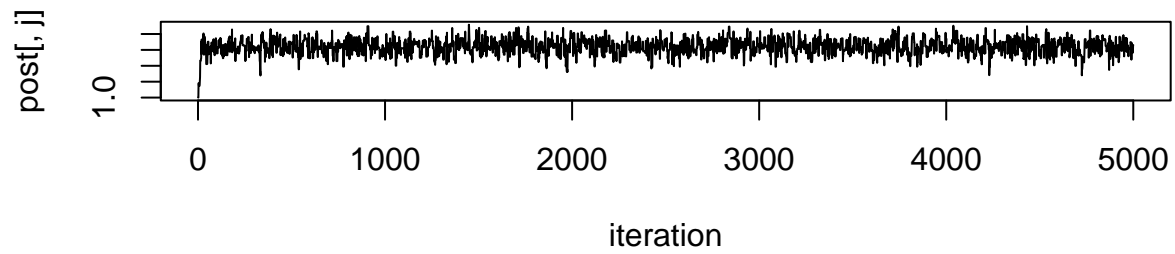
tau











From part b):

```
gradient = function(lamb, tau, beta) {
  nu = 1 / tau
  mu = X %*% beta
  g1 = (lamb - mu) * nu - y + exp(lamb)
  return(as.vector(g1))
}
```

- Here use Leapfrog path

```
leap = function(x0, d0, epsilon, nn, M, tau, beta) {
  Mi = M
  diag(Mi) = 1 / diag(2 * M)
  p = length(x0)
  xx = dd = matrix(0, nrow = p, ncol = nn)
  xx[, 1] = x0
  dd[, 1] = d0
  for(i in 1:(nn - 1)) {
    dd[, (i + 1)] = dd[, i] - 0.5 * epsilon * gradient(xx[, i], tau, beta)
    xx[, (i + 1)] = xx[, i] + epsilon * gradient(dd[, (i + 1)] %% Mi, tau, beta)
    dd[, (i + 1)] = dd[, (i + 1)] - 0.5 * epsilon * gradient(xx[, (i + 1)], tau, beta)
  }
  return(list(xx = xx, dd = dd))
}
```

- HMC sampler

```
hmc = function(y, X, n_iter, burn, eta, L, eps) {

  ## initialize
  n = length(y)
  p = ncol(X)
  XT = t(X)
  XTX = XT %% X
  XTXI = ginv(XTX)
  g = n
  mdelta = rep(0, n)
  M = eta * diag(n)

  ## save chain
  lamb_now = rnorm(n, log(y + 0.01), 0.1)
  beta_now = XTXI %% XT %% lambda
  tau_now = sum(lambda - X %% beta)^2 / (n - p)

  beta_out = matrix(NA, nrow = n_iter, ncol = p)
  lamb_out = matrix(NA, nrow = n_iter, ncol = n)
  tau_out = rep(NA, n_iter)

  ## Sampler
  for(i in 1:n_iter){
    m0 = X %% beta_now
    delta = as.vector(rmvnorm(1, mdelta, M))
    p0 = dmvmnorm(delta, mdelta, M, log = TRUE)
    u0 = sum(-(lamb_now - m0)^2 / (2 * tau_now) + lamb_now * y - exp(lamb_now))

    e1 = rgamma(1, (eps * 200 + 1), rate = 200)
    L1 = max(2, rpois(1, L + 1))
    new = leap(x0 = lamb_now,
              d0 = delta,
              epsilon = e1,
```

```

        nn = L1,
        M,
        tau_now,
        beta_now)

l1 = as.vector(new$xx[,L1])
delta1 = as.vector(new$dd[,L1])

p1 = dmvnorm(delta1, mdelta, M, log = TRUE)
u1 = sum(-(l1 - m0)^2 / (2 * tau) + l1 * y - exp(l1))

ratio = u1 - u0 + p1 - p0

if(log(runif(1)) < ratio) {
  lamb_now = l1
  delta = delta1
}

# beta
beta_now = update_beta(g = g, X = X, lamb = lamb_now, tau = tau_now)

# tau
tau_now = update_tau(n = n, X = X, beta = beta_now, lamb = lamb_now)

if(i > burn){
  i1 = i - burn
  lamb_out[i1,] = lamb_now
  beta_out[i1,] = beta_now
  tau[i1] = tau_now
}

}
cbind(beta = beta_out, lambda = lamb_out, tau = tau_out)
}

```

```

set.seed(888)
post1 = hmc(y, X, n_iter = 5000, burn = 1000, eta = 1, L = 2, eps = 0.01)

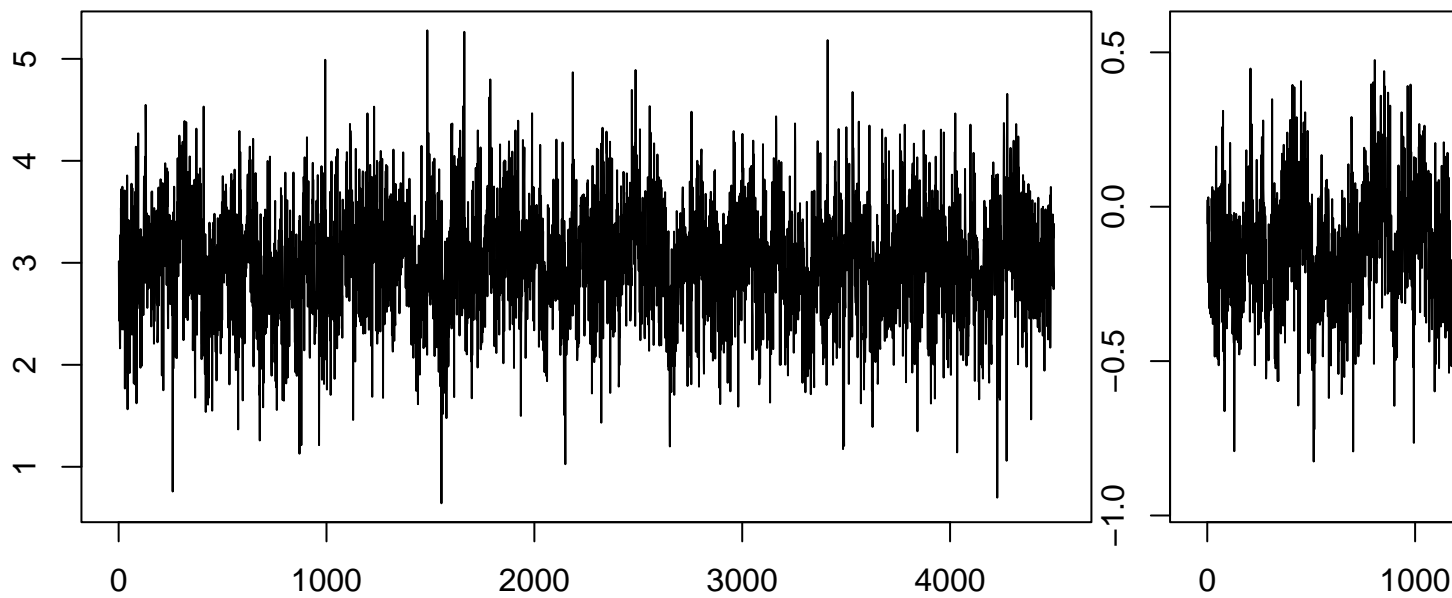
```

```

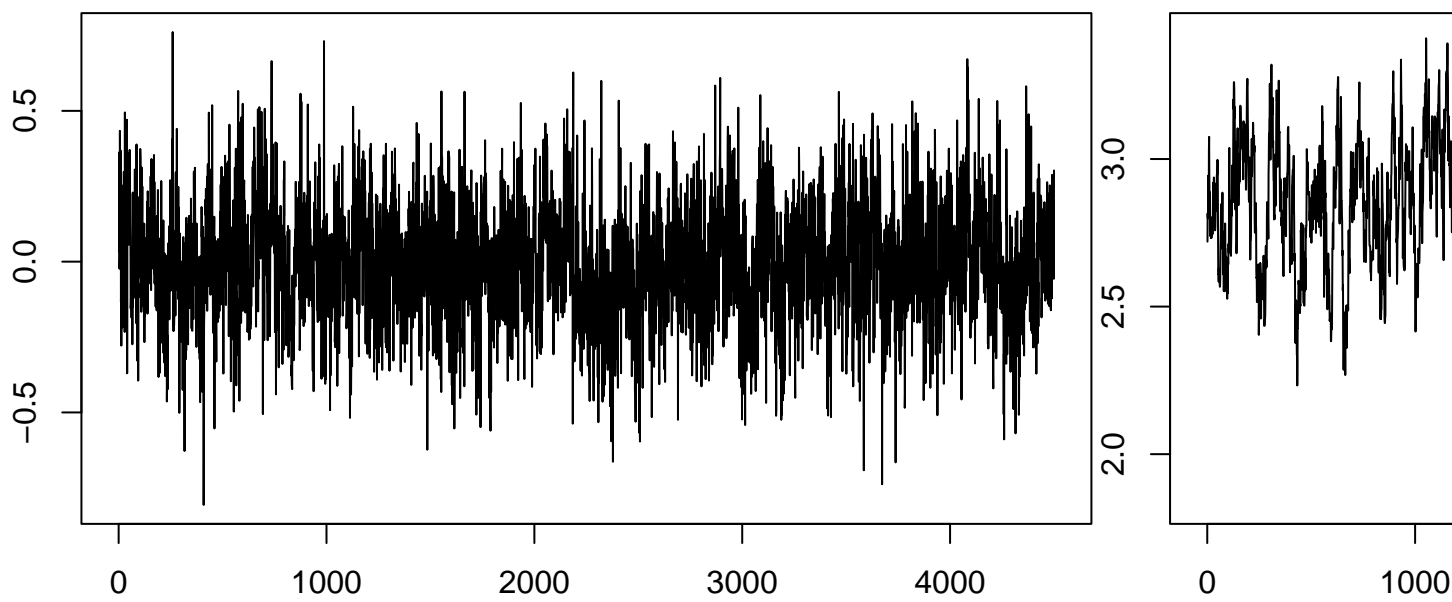
traceplot(as.mcmc(post1))

```

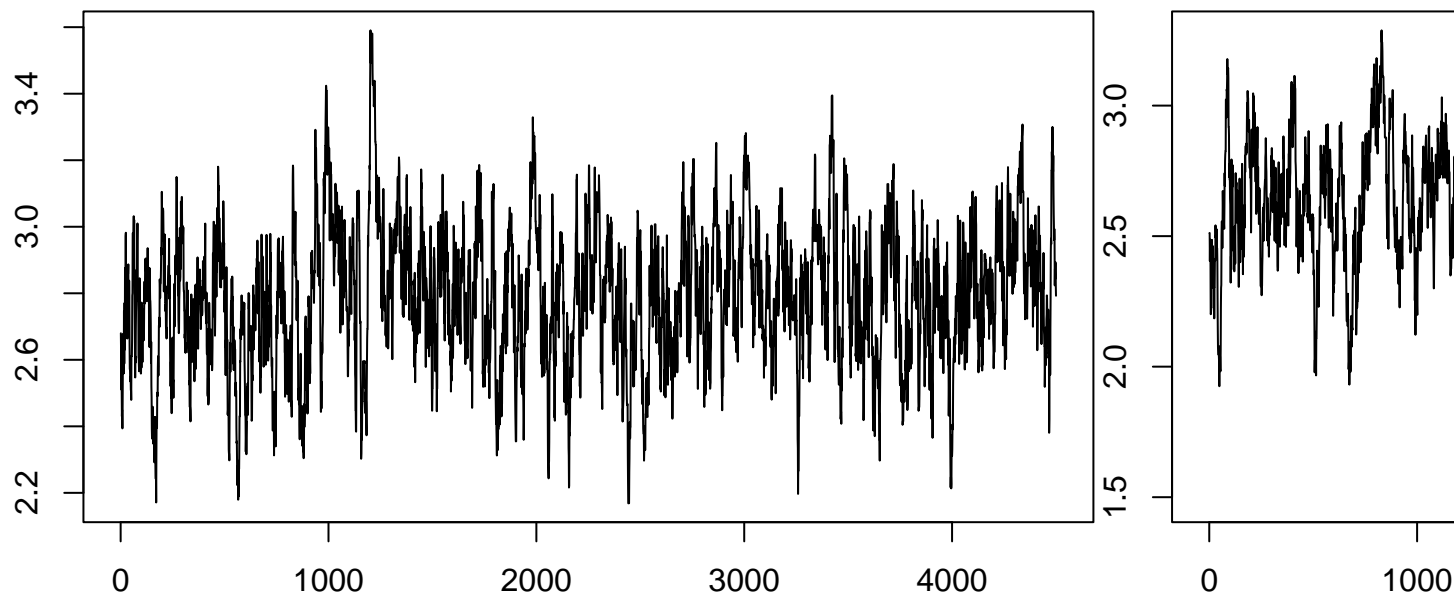
Trace of



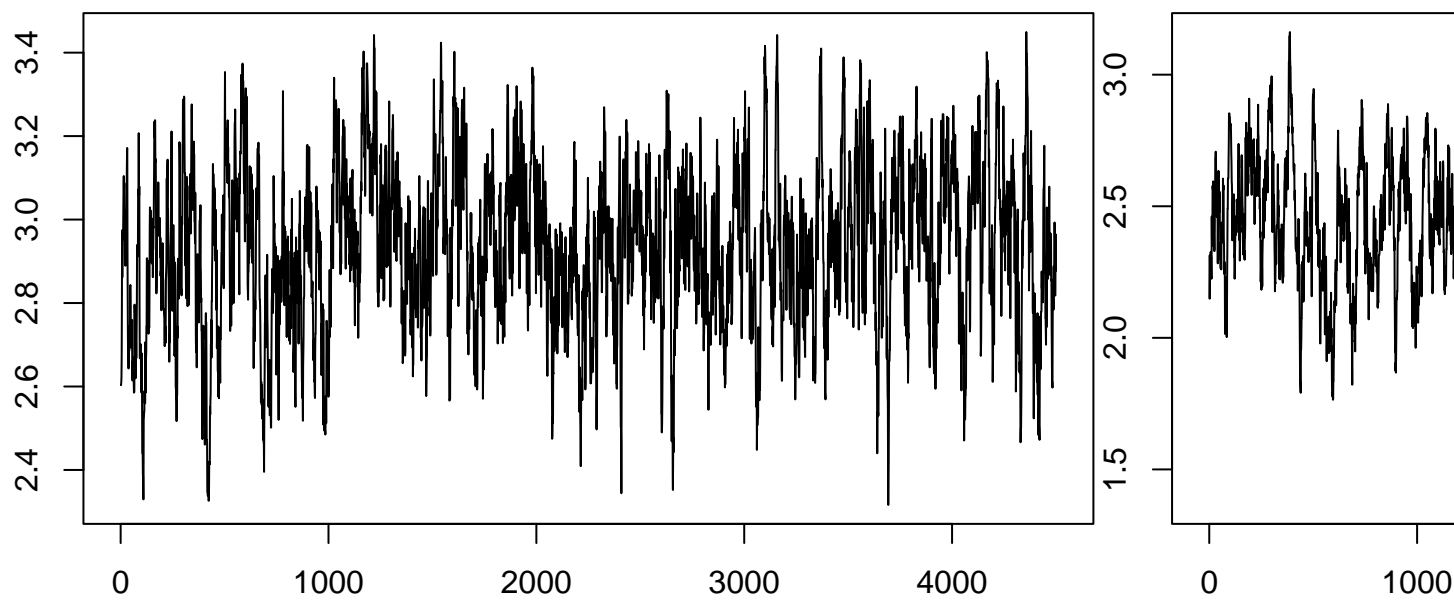
Iterations
Trace of



Trace of

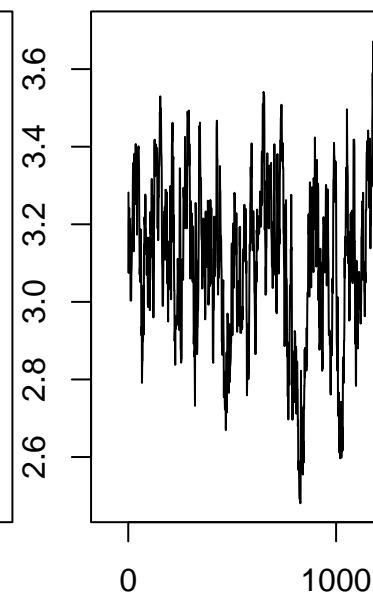
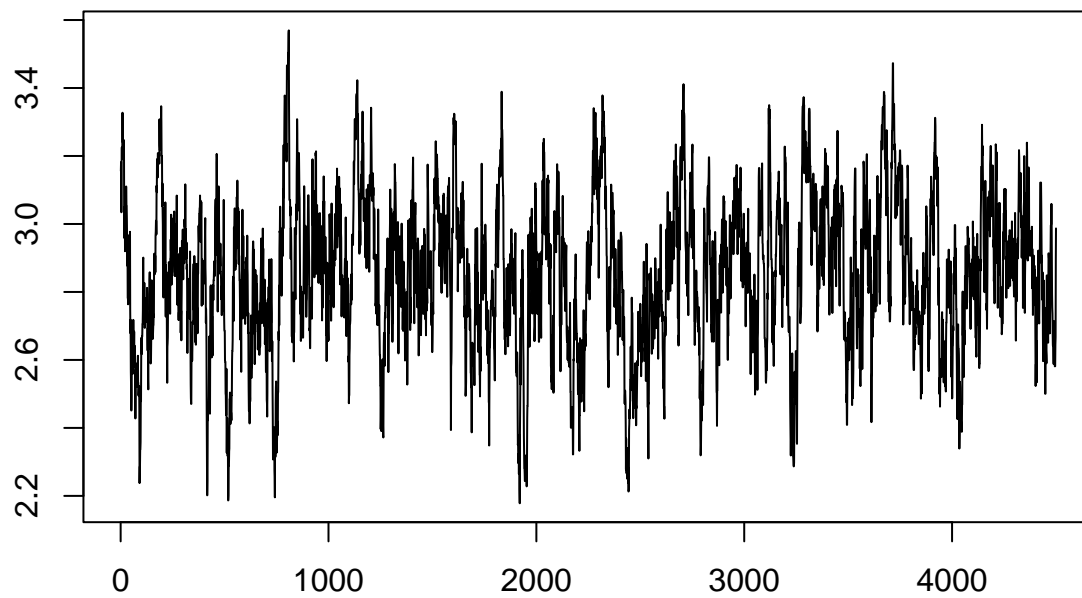


Iterations
Trace of

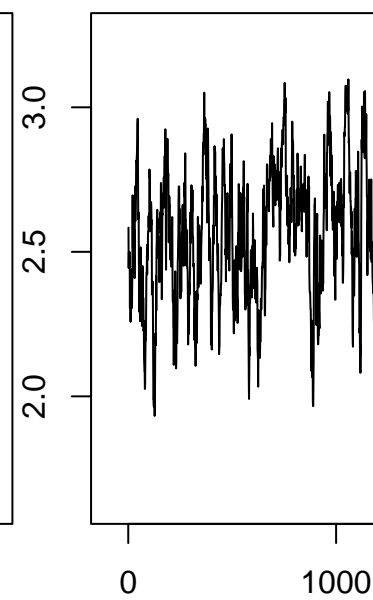
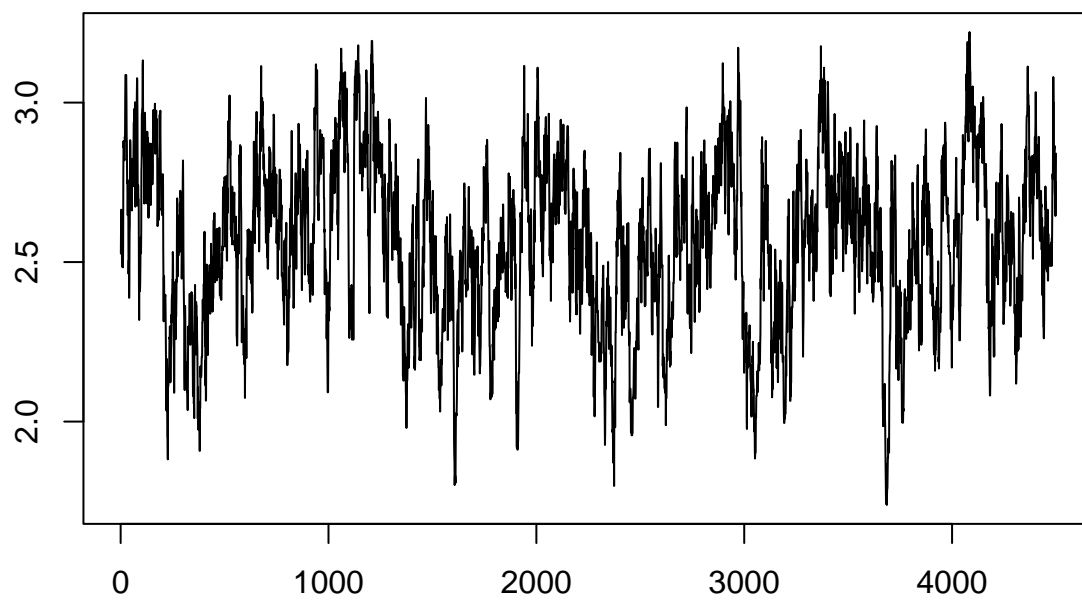


Iterations

Trace of

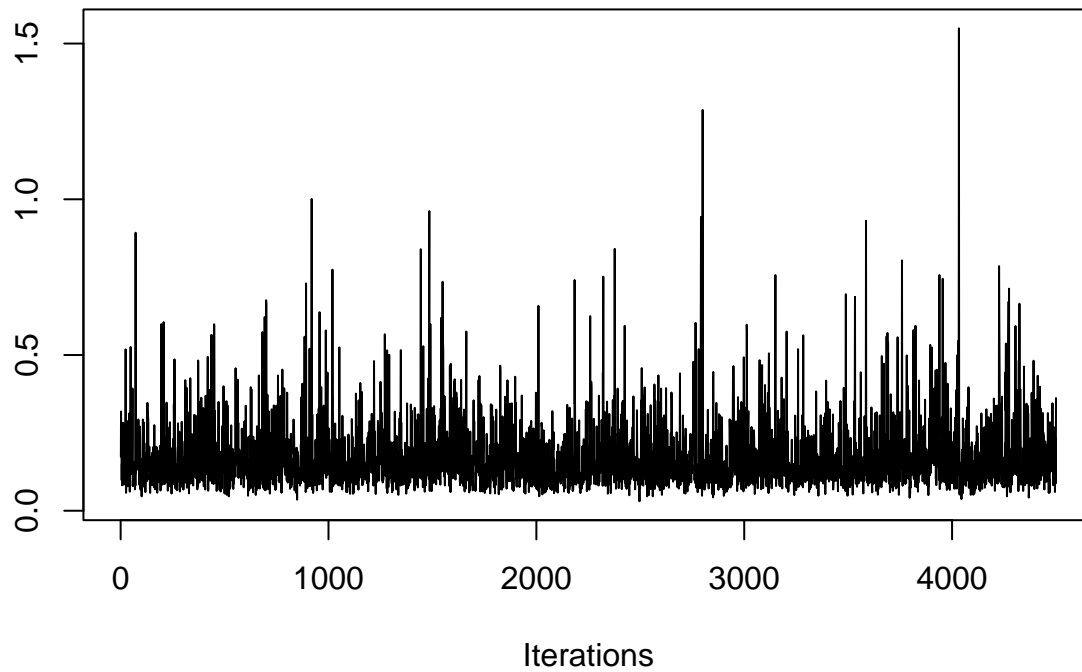


Iterations
Trace of



Iterations

Trace of tau



- All algorithms converged with 5000 simulations and 1,000 burn-in iterations based on the ACF and trace plots. On run-time usage, very similar in both.
- As we can tell from the plots, HMC method improved the parameters convergence and mixing.